

**\*\*MINOR PROJECT\*\***

**\*\*SIGN LANGUAGE TO TEXT\*\***

1.Reedam Ranjan (1805086)  
2.Nandini Kallia (1805075)  
3.Sourdeep Dhar (1805715)

Step 1 — Preparing the Sign Language Classification Dataset(Preprocessing)

```
In [1]: from torch.utils.data import Dataset
from torch.autograd import Variable
import torchvision.transforms as transforms
import torch.nn as nn
import numpy as np
import torch

from typing import List
import csv

class SignLanguageMNIST(Dataset):
    """Sign language classification dataset.

    Utility for loading Sign Language dataset into PyTorch. Dataset posted on
    Kaggle [in 2017, by an unnamed author with username 'tepcerson' :
    https://www.kaggle.com/dataumage/sign-language-mnist

    Each sample is 1 x 1 x 28 x 28, and each label is a scalar.
    """

    @staticmethod
    def get_label_mapping():
        """
        We map all labels to [0, 23]. This mapping from dataset labels [0, 23]
        to letter indices [0, 25] is returned below.
        """
        mapping = list(range(25))
        mapping.pop(9)
        return mapping

    @staticmethod
    def read_label_samples_from_csv(path: str):
        """
        Assumes first column in CSV is the label and subsequent 28x2 values
        are image pixel values 0-255.
        """
        mapping = SignLanguageMNIST.get_label_mapping()
        labels, samples = [], []
        with open(path) as f:
            _ = next(f) # skip header
            for line in csv.reader(f):
                label = int(line[0])
                labels.append(mapping.index(label))
                samples.append(list(map(int, line[1:]))))
        return labels, samples

    def __init__(self,
                 path: str="data/sign_mnist_train.csv",
                 mean: List[float]=[0.486],
                 std: List[float]=[0.229]):
        """
        Args:
            path: Path to '.csv' file containing 'label', 'pixel0', 'pixel1'...
        """
        labels, samples = SignLanguageMNIST.read_label_samples_from_csv(path)
        self.samples = np.array(samples, dtype=np.uint8).reshape((-1, 28, 28, 1))
        self._labels = np.array(labels, dtype=np.uint8).reshape((-1, 1))

        self._mean = mean
        self._std = std

    def __len__(self):
        return len(self._labels)

    def __getitem__(self, idx):
        transform = transforms.Compose([
            transforms.ToPILImage(),
            transforms.RandomResizedCrop(28, scale=(0.8, 1.2)),
            transforms.ToTensor(),
            transforms.Normalize(mean=self._mean, std=self._std)])

        return {
            'image': transform(self.samples[idx]).float(),
            'label': torch.from_numpy(self._labels[idx]).float()}

In [4]: def get_train_test_loaders(batch_size=32):
trainset = SignLanguageMNIST('data/sign_mnist_train.csv')
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True)

testset = SignLanguageMNIST('data/sign_mnist_test.csv')
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False)
return trainloader, testloader

In [5]: loader, _ = get_train_test_loaders(2)
print(next(iter(loader)))

{'image': tensor([[[[0.8961, 0.9132, 0.9132, ..., 0.8961, 0.9132, 0.9132],
[0.9132, 0.9132, 0.9132, ..., 0.9132, 0.9132, 0.9383],
[0.9383, 0.9383, 0.9383, ..., 0.9383, 0.9383, 0.9474],
...,
[1.1187, 1.1015, 1.1708, ..., 1.2557, 1.2043, 1.2043],
[1.1358, 1.1157, 1.1520, ..., 1.2843, 1.2043, 1.1972],
[1.1358, 1.1358, 1.1708, ..., 1.2843, 1.2843, 1.2214]]]],
[0.6221, 0.6392, 0.6563, ..., 0.4337, 0.4166, 0.3994],
[0.6906, 0.6906, 0.6906, ..., 0.4851, 0.4851, 0.4588],
[0.7248, 0.7248, 0.7248, ..., 0.5364, 0.5193, 0.5022],
...,
[1.2557, 1.2395, 1.2214, ..., 1.1015, 1.0844, 1.0073],
[1.2557, 1.2557, 1.2557, ..., 1.1187, 1.1015, 1.0073],
[1.2557, 1.2395, 1.2557, ..., 1.1187, 1.0844, 1.0073]]]], 'label': tensor([11.],
[13.])])

Step 2 — Building and Training the Sign Language Classifier Using Deep Learning
```

```
In [6]: from torch.utils.data import Dataset
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 3)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 6, 3)
        self.conv3 = nn.Conv2d(6, 16, 3)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 48)
        self.fc3 = nn.Linear(48, 25)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

def main():
    net = Net().float()
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.0)
    scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

    trainloader, _ = get_train_test_loaders()
    for epoch in range(12): # loop over the dataset multiple times
        train(net, criterion, optimizer, trainloader, epoch)
        scheduler.step()
        torch.save(net.state_dict(), "checkpoint.pth")

    def train(net, criterion, optimizer, trainloader, epoch):
        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            inputs = Variable(data['image']).float()
            labels = Variable(data['label']).long()
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels[:, 0])
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            if i % 100 == 0:
                print('[%d, %5d] loss: %.6f' % (epoch, i, running_loss / (i + 1)))

In [7]: main()

[0, 0] loss: 3.193408
[0, 100] loss: 3.209208
[0, 200] loss: 3.209208
[0, 300] loss: 3.020502
[0, 400] loss: 2.767063
[0, 500] loss: 2.494406
[0, 600] loss: 2.243822
[0, 700] loss: 2.038367
[0, 800] loss: 1.863631
[1, 0] loss: 0.660421
[1, 100] loss: 0.454788
[1, 200] loss: 0.431698
[1, 300] loss: 0.401833
[1, 400] loss: 0.397726
[1, 500] loss: 0.342959
[1, 600] loss: 0.325167
[1, 700] loss: 0.396999
[1, 800] loss: 0.291484
[2, 0] loss: 0.947885
[2, 100] loss: 0.160940
[2, 200] loss: 0.149998
[2, 300] loss: 0.146902
[2, 400] loss: 0.142737
[2, 500] loss: 0.133513
[2, 600] loss: 0.122887
[2, 700] loss: 0.130138
[2, 800] loss: 0.125888
[3, 0] loss: 0.031571
[3, 100] loss: 0.090978
[3, 200] loss: 0.084411
[3, 300] loss: 0.081451
[3, 400] loss: 0.080571
[3, 500] loss: 0.084438
[3, 600] loss: 0.080302
[3, 700] loss: 0.087649
[3, 800] loss: 0.087946
[4, 0] loss: 0.047777
[4, 100] loss: 0.057266
[4, 200] loss: 0.061120
[4, 300] loss: 0.060221
[4, 400] loss: 0.058293
[4, 500] loss: 0.059030
[4, 600] loss: 0.058506
[4, 700] loss: 0.058511
[4, 800] loss: 0.057770
[5, 0] loss: 0.007457
[5, 100] loss: 0.008351
[5, 200] loss: 0.008709
[5, 300] loss: 0.005211
[5, 400] loss: 0.008836
[5, 500] loss: 0.008849
[5, 600] loss: 0.007620
[5, 700] loss: 0.005358
[5, 800] loss: 0.005398
[6, 0] loss: 0.072571
[6, 100] loss: 0.032321
[6, 200] loss: 0.048233
[6, 300] loss: 0.036612
[6, 400] loss: 0.036078
[6, 500] loss: 0.041315
[6, 600] loss: 0.042208
[6, 700] loss: 0.040826
[6, 800] loss: 0.040325
[7, 0] loss: 0.008301
[7, 100] loss: 0.030616
[7, 200] loss: 0.031134
[7, 300] loss: 0.030806
[7, 400] loss: 0.032173
[7, 500] loss: 0.033007
[7, 600] loss: 0.030827
[7, 700] loss: 0.030285
[7, 800] loss: 0.032723
[8, 0] loss: 0.150842
[8, 100] loss: 0.041839
[8, 200] loss: 0.041024
[8, 300] loss: 0.039950
[8, 400] loss: 0.045207
[8, 500] loss: 0.046381
[8, 600] loss: 0.043453
[8, 700] loss: 0.040535
[8, 800] loss: 0.039876
[9, 0] loss: 0.003689
[9, 100] loss: 0.030894
[9, 200] loss: 0.030772
[9, 300] loss: 0.030917
[9, 400] loss: 0.034002
[9, 500] loss: 0.030806
[9, 600] loss: 0.032566
[9, 700] loss: 0.032744
[9, 800] loss: 0.031113
[10, 0] loss: 0.000153
[10, 100] loss: 0.017307
[10, 200] loss: 0.010505
[10, 300] loss: 0.013979
[10, 400] loss: 0.012066
[10, 500] loss: 0.011855
[10, 600] loss: 0.011260
[10, 700] loss: 0.010506
[10, 800] loss: 0.010019
[11, 0] loss: 0.000162
[11, 100] loss: 0.000919
[11, 200] loss: 0.006931
[11, 300] loss: 0.005598
[11, 400] loss: 0.007027
[11, 500] loss: 0.006762
[11, 600] loss: 0.006077
[11, 700] loss: 0.005940
[11, 800] loss: 0.005476

In [8]: main()

[0, 0] loss: 3.218370
[0, 100] loss: 3.209821
[0, 200] loss: 3.202712
[0, 300] loss: 3.133660
[0, 400] loss: 2.929101
[0, 500] loss: 2.690807
[0, 600] loss: 2.455835
[0, 700] loss: 2.228534
[0, 800] loss: 2.027808
[1, 0] loss: 0.507166
[1, 100] loss: 0.500027
[1, 200] loss: 0.461239
[1, 300] loss: 0.417974
[1, 400] loss: 0.390625
[1, 500] loss: 0.369000
[1, 600] loss: 0.344547
[1, 700] loss: 0.321961
[1, 800] loss: 0.304202
[2, 0] loss: 0.435642
[2, 100] loss: 0.179268
[2, 200] loss: 0.153568
[2, 300] loss: 0.157750
[2, 400] loss: 0.152462
[2, 500] loss: 0.143553
[2, 600] loss: 0.142807
[2, 700] loss: 0.138218
[2, 800] loss: 0.135241
[3, 0] loss: 0.203142
[3, 100] loss: 0.093122
[3, 200] loss: 0.092518
[3, 300] loss: 0.089966
[3, 400] loss: 0.080251
[3, 500] loss: 0.080301
[3, 600] loss: 0.080322
[3, 700] loss: 0.085312
[3, 800] loss: 0.085651
[4, 0] loss: 0.039505
[4, 100] loss: 0.005106
[4, 200] loss: 0.075207
[4, 300] loss: 0.060901
[4, 400] loss: 0.069145
[4, 500] loss: 0.065298
[4, 600] loss: 0.065075
[4, 700] loss: 0.065393
[4, 800] loss: 0.069289
[5, 0] loss: 0.030225
[5, 100] loss: 0.050000
[5, 200] loss: 0.046112
[5, 300] loss: 0.047215
[5, 400] loss: 0.046595
[5, 500] loss: 0.043227
[5, 600] loss: 0.043054
[5, 700] loss: 0.044323
[5, 800] loss: 0.044103
[6, 0] loss: 0.061230
[6, 100] loss: 0.031348
[6, 200] loss: 0.040579
[6, 300] loss: 0.040135
[6, 400] loss: 0.046402
[6, 500] loss: 0.044303
[6, 600] loss: 0.044723
[6, 700] loss: 0.041663
[6, 800] loss: 0.043448
[7, 0] loss: 0.021581
[7, 100] loss: 0.005508
[7, 200] loss: 0.043741
[7, 300] loss: 0.039236
[7, 400] loss: 0.038254
[7, 500] loss: 0.038604
[7, 600] loss: 0.038711
[7, 700] loss: 0.037918
[7, 800] loss: 0.035991
[8, 0] loss: 0.002241
[8, 100] loss: 0.030937
[8, 200] loss: 0.032439
[8, 300] loss: 0.032727
[8, 400] loss: 0.034107
[8, 500] loss: 0.035072
[8, 600] loss: 0.033995
[8, 700] loss: 0.032971
[8, 800] loss: 0.033962
[9, 0] loss: 0.004006
[9, 100] loss: 0.034063
[9, 200] loss: 0.031797
[9, 300] loss: 0.031906
[9, 400] loss: 0.031440
[9, 500] loss: 0.033291
[9, 600] loss: 0.033475
[9, 700] loss: 0.031474
[9, 800] loss: 0.031825
[10, 0] loss: 0.002000
[10, 100] loss: 0.010861
[10, 200] loss: 0.011170
[10, 300] loss: 0.012243
[10, 400] loss: 0.012188
[10, 500] loss: 0.011350
[10, 600] loss: 0.011152
[10, 700] loss: 0.010790
[10, 800] loss: 0.010062
[11, 0] loss: 0.008281
[11, 100] loss: 0.004162
[11, 200] loss: 0.005785
[11, 300] loss: 0.006706
[11, 400] loss: 0.006712
[11, 500] loss: 0.006371
[11, 600] loss: 0.006017
[11, 700] loss: 0.005722
[11, 800] loss: 0.005538

Step 3 — Evaluating the Sign Language Classifier
```

```
In [9]: from torch.utils.data import Dataset
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch
import numpy as np

import onnx
import onnxruntime as ort

def evaluate(outputs: Variable, labels: Variable) -> float:
    """Evaluate neural network outputs against non-one-hotthed labels."""
    y = labels.numpy()
    yhat = np.argmax(outputs, axis=1)
    return float(np.sum(yhat == y))

def batch_evaluate(
    net: Net,
    dataloader: torch.utils.data.DataLoader) -> float:
    """Evaluate neural network in batches, if dataset is too large."""
    score = 0
    for batch in dataloader:
        n = len(batch['image'])
        outputs = net(batch['image'])
        if isinstance(outputs, torch.Tensor):
            outputs = outputs.detach().numpy()
        score = evaluate(outputs, batch['label'][:, 0])
    return score

def validate():
    trainloader, testloader = get_train_test_loaders()
    net = Net().float().eval()

    pretrained_model = torch.load("checkpoint.pth")
    net.load_state_dict(pretrained_model)

    print('=' * 10, 'PyTorch', '=' * 10)
    train_acc = batch_evaluate(net, trainloader) * 100.
    print('Training accuracy: %.1f' % train_acc)
    test_acc = batch_evaluate(net, testloader) * 100.
    print('Validation accuracy: %.1f' % test_acc)

    trainloader, testloader = get_train_test_loaders(1)

    # export to onnx
    fname = "signlanguage.onnx"
    dummy = torch.randn(1, 1, 28, 28)
    torch.onnx.export(net, dummy, fname, input_names=['input'])

    # check exported model
    model = onnx.load(fname)
    onnx.checker.check_model(model) # check model is well-formed

    # create runnable session with exported model
    ort_session = ort.InferenceSession(fname)
    net = LambdaLmp: ort_session.run(None, ('input': inp.data.numpy()))[0]

    print('=' * 10, 'ONNX', '=' * 10)
    train_acc = batch_evaluate(net, trainloader) * 100.
    print('Training accuracy: %.1f' % train_acc)
    test_acc = batch_evaluate(net, testloader) * 100.
    print('Validation accuracy: %.1f' % test_acc)

In [10]: validate()

===== PyTorch =====
Training accuracy: 99.8
Validation accuracy: 97.2
===== ONNX =====
Training accuracy: 99.8
Validation accuracy: 96.9

Step 4 — Linking the Camera Feed
```

```
In [11]: import cv2
import numpy as np
import onnxruntime as ort

def center_crop(frame):
    h, w = frame.shape
    start = abs(h - w) // 2
    if h > w:
        return frame[start: start + w]
    return frame[:, start: start + h]

def main_ca():
    # constants
    index_to_letter = list('ABCDEFGHIKLMNOPQRSTUVWXYZ')
    mean = 0.485 * 255.
    std = 0.229 * 255.

    # create runnable session with exported model
    ort_session = ort.InferenceSession("signlanguage.onnx")

    cap = cv2.VideoCapture(0)
    while True:
        # capture frame-by-frame
        ret, frame = cap.read()

        # preprocess data
        frame = center_crop(frame)
        frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
        x = cv2.resize(frame, (28, 28))
        x = (x - mean) / std

        x = x.reshape(1, 1, 28, 28).astype(np.float32)
        y = ort_session.run(None, ('input': x))[0]

        index = np.argmax(y, axis=1)
        letter = index_to_letter[int(index)]

        cv2.putText(frame, letter, (100, 100), cv2.FONT_HERSHEY_SIMPLEX, 2.0, (0, 255, 0), thickness=2)
        cv2.imshow("Sign Language Translator", frame)

        if cv2.waitKey(1) & 0xFF == ord('q'):
            break

    cap.release()
    cv2.destroyAllWindows()

In [ ]: main_ca()
```